

## Cheat Sheet: AI Models for NLP

Package/Method	Description	Code example
PyTorch/Embedding and EmbeddingBag	<p>Embedding is a class that represents an embedding layer. It accepts token indices and produces embedding vectors.</p> <p>EmbeddingBag is a class that aggregates embeddings using mean or sum operations.</p> <p>Embedding and EmbeddingBag are part of the torch.nn module. The code example shows how you can use Embedding and EmbeddingBag in PyTorch.</p>	<pre># Defining a data set dataset = [     "I like cats",     "I hate dogs",     "I'm impartial to hippos" ] #Initializing the tokenizer, iterator from the data set, and vocabulary tokenizer = get_tokenizer('spacy', language='en_core_web_sm') def yield_tokens(data_iter):     for data_sample in data_iter:         yield tokenizer(data_sample) data_iter = iter(dataset) vocab = build_vocab_from_iterator(yield_tokens(data_iter)) #Tokenizing and generating indices input_ids=lambda x:[torch.tensor(vocab(tokenizer(data_sample))) for data_sample in dataset] index=input_ids(dataset) print(index) #Initiating the embedding layer, specifying the dimension size for the embeddings, #determining the count of unique tokens present in the vocabulary, and creating the embedding layer embedding_dim = 3 n_embedding = len(vocab) n_embedding:9 embeds = nn.Embedding(n_embedding, embedding_dim) #Applying the embedding object i_like_cats=embeds(index[0]) i_like_cats impartial_to_hippos=embeds(index[-1]) impartial_to_hippos #Initializing the embedding bag layer embedding_dim = 3 n_embedding = len(vocab) n_embedding:9 embedding_bag = nn.EmbeddingBag(n_embedding, embedding_dim) # Output the embedding bag dataset = ["I like cats","I hate dogs","I'm impartial to hippos"] index:[tensor([0, 7, 2]), tensor([0, 4, 3]), tensor([0, 1, 6, 8, 5])] i_like_cats=embedding_bag(index[0],offsets=torch.tensor([0])) i_like_cats</pre>
Batch function	<p>Defines the number of samples that will be propagated through the network.</p>	<pre>def collate_batch(batch):     target_list, context_list, offsets = [], [], [0]     for _context, _target in batch:         target_list.append(vocab[_target])         processed_context = torch.tensor(text_pipeline(_context), dtype=torch.int64)         context_list.append(processed_context)         offsets.append(processed_context.size(0))     target_list = torch.tensor(target_list, dtype=torch.int64)     offsets = torch.tensor(offsets[:-1]).cumsum(dim=0)     context_list = torch.cat(context_list)     return target_list.to(device), context_list.to(device), offsets.to(device) BATCH_SIZE = 64 # batch size for training dataloader_cbow = DataLoader(cbow_data, batch_size=BATCH_SIZE, shuffle=True, collate_fn=collate_batch)</pre>

Package/Method	Description	Code example
Forward pass	Refers to the computation and storage of intermediate variables (including outputs) for a neural network in order from the input to the output layer.	<pre>def forward(self, text):</pre>
Stanford's pre-trained GloVe	Leverages large-scale data for word embeddings. It can be integrated into PyTorch for improved NLP tasks such as classification.	<pre>from torchtext.vocab import GloVe, vocab # Creating an instance of the 6B version of GloVe() model glove_vectors_6B = GloVe(name='6B') # you can specify the model with the following format: GloVe(name='6B') # Build vocab from glove_vectors vocab = vocab(glove_vectors_6B.stoi, 0, specials=('&lt;unk&gt;', '&lt;pad&gt;')) vocab.set_default_index(vocab["&lt;unk&gt;"])</pre>
vocab	The vocab object is part of the PyTorch torchtext library. It maps tokens to indices. The code example shows how you can apply the vocab object to tokens directly.	<pre># Takes an iterator as input and extracts the next tokenized sentence. Creates a list of token indices def get_tokenized_sentence_and_indices(iterator):     tokenized_sentence = next(iterator)     token_indices = [vocab[token] for token in tokenized_sentence]     return tokenized_sentence, token_indices  # Returns the tokenized sentences and the corresponding token indices. Repeats the process. tokenized_sentence, token_indices = get_tokenized_sentence_and_indices(my_iterator) next(my_iterator) # Prints the tokenized sentence and its corresponding token indices. print("Tokenized Sentence:", tokenized_sentence) print("Token Indices:", token_indices)</pre>
Special tokens in PyTorch: <eos> and <bos>	Tokens introduced to input sequences to convey specific information or serve a particular purpose during training. The code example shows the use of <bos> and <eos> during tokenization. The <bos> token denotes the beginning of the input sequence, and the <eos> token denotes the end.	<pre># Appends &lt;bos&gt; at the beginning and &lt;eos&gt; at the end of the tokenized sentences # using a loop that iterates over the sentences in the input data tokenizer_en = get_tokenizer('spacy', language='en_core_web_sm') tokens = [] max_length = 0 for line in lines:     tokenized_line = tokenizer_en(line)     tokenized_line = ['&lt;bos&gt;'] + tokenized_line + ['&lt;eos&gt;']     tokens.append(tokenized_line)     max_length = max(max_length, len(tokenized_line))</pre>
Special tokens in PyTorch: <pad>	Tokens introduced to input sequences	<pre># Pads the tokenized lines for i in range(len(tokens)):     tokens[i] = tokens[i] + [&lt;pad&gt;] * (max_length - len(tokens[i]))</pre>

Package/Method	Description	Code example
	to convey specific information or serve a particular purpose during training. The code example shows the use of <pad> token to ensure all sentences have the same length.	
Cross entropy loss	A metric used in machine learning (ML) to evaluate the performance of a classification model. The loss is measured as the probability value between 0 (perfect model) and 1. Typically, the aim is to bring the model as close to 0 as possible.	<pre>from torch.nn import CrossEntropyLoss model = TextClassificationModel(vocab_size, emsize, num_class) loss_fn = CrossEntropyLoss() predicted_label = model(text, offsets) loss = criterion(predicted_label, label)</pre>
Optimization	Method to reduce losses in a model.	<pre># Creates an iterator object optimizer = torch.optim.SGD(model.parameters(), lr=0.1) scheduler = torch.optim.lr_scheduler.StepLR(optimizer, 1.0, gamma=0.1) optimizer.zero_grad() predicted_label = model(text, offsets) loss = criterion(predicted_label, label) loss.backward() torch.nn.utils.clip_grad_norm_(model.parameters(), 0.1) optimizer.step()</pre>
sentence_bleu()	NLTK (or Natural Language Toolkit) provides this function to evaluate a hypothesis sentence against one or more reference sentences. The reference sentences must be presented as a list of sentences where each reference is a list of tokens.	<pre>from nltk.translate.bleu_score import sentence_bleu def calculate_bleu_score(generated_translation, reference_translations):     # Convert the generated translations and reference translations into the expected format for sentence_bleu     references = [reference.split() for reference in reference_translations]     hypothesis = generated_translation.split()     # Calculate the BLEU score     bleu_score = sentence_bleu(references, hypothesis)     return bleu_score reference_translations = ["Asian man sweeping the walkway .", "An asian man sweeping the walkway .", "A man sweeping the walkway ."] bleu_score = calculate_bleu_score(generated_translation, reference_translations)</pre>
Encoder RNN model	The encoder-decoder seq2seq model works together to transform an	<pre>class Encoder(nn.Module):     def __init__(self, vocab_len, emb_dim, hid_dim, n_layers, dropout_prob):         super().__init__()         self.hid_dim = hid_dim         self.n_layers = n_layers</pre>

Package/Method	Description	Code example
	input sequence into an output sequence. Encoder is a series of RNNs that process the input sequence individually, passing their hidden states to their next RNN.	<pre>self.embedding = nn.Embedding(vocab_len, emb_dim) self.lstm = nn.LSTM(emb_dim, hid_dim, n_layers, dropout = dropout_prob) self.dropout = nn.Dropout(dropout_prob) def forward(self, input_batch):     embed = self.dropout(self.embedding(input_batch))     embed = embed.to(device)     outputs, (hidden, cell) = self.lstm(embed)     return hidden, cell</pre>
Decoder RNN model	The encoder-decoder seq2seq model works together to transform an input sequence into an output sequence. The decoder module is a series of RNNs that autoregressively generates the translation as one token at a time. Each generated token goes back into the next RNN along with the hidden state to generate the next token of the output sequence until the end token is generated.	<pre>class Decoder(nn.Module):     def __init__(self, output_dim, emb_dim, hid_dim, n_layers, dropout):         super().__init__()         self.output_dim = output_dim         self.hid_dim = hid_dim         self.n_layers = n_layers         self.embedding = nn.Embedding(output_dim, emb_dim)         self.lstm = nn.LSTM(emb_dim, hid_dim, n_layers, dropout = dropout)         self.fc_out = nn.Linear(hid_dim, output_dim)         self.softmax = nn.LogSoftmax(dim=1)         self.dropout = nn.Dropout(dropout)     def forward(self, input, hidden, cell):         input = input.unsqueeze(0)         embedded = self.dropout(self.embedding(input))         output, (hidden, cell) = self.lstm(embedded, (hidden, cell))         prediction_logit = self.fc_out(output.squeeze(0))         prediction = self.softmax(prediction_logit)         return prediction, hidden, cell</pre>
Skip-gram model	Predicts surrounding context words from a specific target word. It predicts one context word at a time from a target word.	<pre>class SkipGram_Model(nn.Module):     def __init__(self, vocab_size, embed_dim):         super(SkipGram_Model, self).__init__()         # Define the embeddings layer         self.embeddings = nn.Embedding(num_embeddings=vocab_size, embedding_dim=embed_dim)         # Define the fully connected layer         self.fc = nn.Linear(in_features=embed_dim, out_features=vocab_size)         # Perform the forward pass     def forward(self, text):         # Pass the input text through the embeddings layer         out = self.embeddings(text)         # Pass the output of the embeddings layer through the fully connected layer         # Apply the ReLU activation function         out = torch.relu(out)         out = self.fc(out)         return out model_sg = SkipGram_Model(vocab_size, emsize).to(device) # Sequence generation function CONTEXT_SIZE = 2 skip_data = [] for i in range(CONTEXT_SIZE, len(tokenized_toy_data) - CONTEXT_SIZE):     context = (         [tokenized_toy_data[i - j - 1] for j in range(CONTEXT_SIZE)] # Preceding words         + [tokenized_toy_data[i + j + 1] for j in range(CONTEXT_SIZE)] # Succeeding words     )     target = tokenized_toy_data[i]     skip_data.append((target, context)) skip_data=[('i', ['wish', 'i', 'was', 'little']), ('was', ['i', 'wish', 'little', 'bit'])],...</pre>

Package/Method	Description	Code example
collate_fn	Processes the list of samples to form a batch. The batch argument is a list of all your samples.	<pre>def collate_fn(batch):     target_list, context_list = [], []     for _context, _target in batch:         target_list.append(vocab[_target])         context_list.append(vocab[_context])     target_list = torch.tensor(target_list, dtype=torch.int64)     context_list = torch.tensor(context_list, dtype=torch.int64)     return target_list.to(device), context_list.to(device)</pre>
Training function	Trains the model for a specified number of epochs. It also includes a condition to check whether the input is for skip-gram or CBOW. The output of this function includes the trained model and a list of average losses for each epoch.	<pre>def train_model(model, dataloader, criterion, optimizer, num_epochs=1000):     # List to store running loss for each epoch     epoch_losses = []     for epoch in tqdm(range(num_epochs)):         # Storing running loss values for the current epoch         running_loss = 0.0         # Using tqdm for a progress bar         for idx, samples in enumerate(dataloader):             optimizer.zero_grad()             # Check for EmbeddingBag layer in the model CBOW             if any(isinstance(module, nn.EmbeddingBag) for _, module in model.named_modules()):                 target, context, offsets = samples                 predicted = model(context, offsets)             # Check for Embedding layer in the model skip gram             elif any(isinstance(module, nn.Embedding) for _, module in model.named_modules()):                 target, context = samples                 predicted = model(context)             loss = criterion(predicted, target)             loss.backward()             torch.nn.utils.clip_grad_norm_(model.parameters(), 0.1)             optimizer.step()             running_loss += loss.item()         # Append average loss for the epoch         epoch_losses.append(running_loss / len(dataloader))     return model, epoch_losses</pre>
CBOW model	Utilizes context words to predict a target word and generate its embedding.	<pre>class CBOW(nn.Module):     # Initialize the CBOW model     def __init__(self, vocab_size, embed_dim, num_class):         super(CBOW, self).__init__()         # Define the embedding layer using nn.EmbeddingBag         self.embedding = nn.EmbeddingBag(vocab_size, embed_dim, sparse=False)         # Define the fully connected layer         self.fc = nn.Linear(embed_dim, vocab_size)     def forward(self, text, offsets):         # Pass the input text and offsets through the embedding layer         out = self.embedding(text, offsets)         # Apply the ReLU activation function to the output of the first linear layer         out = torch.relu(out)         # Pass the output of the ReLU activation through the fully connected layer</pre>

Package/Method	Description	Code example
		<pre>return self.fc(out) vocab_size = len(vocab) emsize = 24 model_cbow = CBOW(vocab_size, emsize, vocab_size).to(device)</pre>
Training loop	Enumerates data from the DataLoader and, on each pass of the loop, gets a batch of training data from the DataLoader, zeros the optimizer's gradients, and performs an inference (gets predictions from the model for an input batch).	<pre>for epoch in tqdm(range(1, EPOCHS + 1)):     model.train()     cum_loss=0     for idx, (label, text, offsets) in enumerate(train_dataloader):         optimizer.zero_grad()         predicted_label = model(text, offsets)         loss = criterion(predicted_label, label)         loss.backward()         torch.nn.utils.clip_grad_norm_(model.parameters(), 0.1)         optimizer.step()         cum_loss+=loss.item()     cum_loss_list.append(cum_loss)     accu_val = evaluate(valid_dataloader)     acc_epoch.append(accu_val)     if accu_val &gt; acc_old:         acc_old= accu_val         torch.save(model.state_dict(), 'my_model.pth')</pre>



**Skills Network**